

Welcome to

# Big Data & Hadoop

Session

---

Session 5 - Pig & Pig Latin



**+91-9538998962**

**sandeep@knowbigdata.com**



---

# WELCOME - KNOWBIGDATA

---

- Interact - Ask Questions
- Real Life Project
- Lifetime access of content
- Quizzes & Certification Test
- Class Recording
- 10 x (3hr class)
- Cluster Access
- Socio-Pro Visibility
- 24x7 support
- Mock Interviews



# ABOUT ME

2014	<b>KnowBigData</b>	Founded
2014	<b>Amazon</b>	Built High Throughput Systems for <a href="http://Amazon.com">Amazon.com</a> site using in-house NoSql.
2012		
2012	<b>InMobi</b>	Built Recommender after churning 200 TB
2011	<b>tBits Global</b>	Founded tBits Global Built an enterprise grade Document Management System
2006	<b>D.E.Shaw</b>	Built the big data systems before the term was coined
2002	<b>IIT Roorkee</b>	Finished B.Tech somehow.
2002		





# COURSE CONTENT

I	Understanding BigData, Hadoop Architecture
II	Environment Overview, MapReduce Basics
III	Adv MapReduce & Testing
➔ IV	Pig & Pig Latin
V	Analytics using Hive
VI	NoSQL, HBASE
VII	Oozie, Mahout,
VIII	Zookeeper, Apache Storm
IX	Apache Flume, Apache Spark
X	YARN, Big Data Sets & Project Assignment



---

# TODAY'S CLASS

---

- Introduction
- Wordcount example
- Pig Latin vs SQL
- Pig Latin vs MapReduce
- Use Cases
- Philosophy
- Installation
- Using from Cloud Labs
- Local Mode
- Command Line Args
- Prompt - Grunt
- Average Val Example
- Data Types - Scaler
- Data Types - Complex
- Schema
- Basic Contructions
- LOAD
- Diagnostic
- STORE & DUMP
- Foreach



---

# TODAY'S CLASS

---

- FILTER
- Order by
- JOIN
- Limit & Sample
- Parallel
- Flatten
- Calling External Functions
- Foreach Nested
- Fragmented Joins
- Skewed Join - Merge
- Left Outer Join
- Right Outer Join
- Full Join
- CoGroup
- Union
- Cross
- Stream
- Params
- Splits
- Non-Linear Execution



---

# ADV. MAP / REDUCE

---

**Followup Questions?**



---

# PIG

---

An engine for executing data flows in parallel on Hadoop

- Language - Pig Latin
  - Provides operators: Join, order by, Filter ...
- Runs on Hadoop. Makes use of both
  - HDFS
  - MapReduce



# PIG - WORD COUNT EXAMPLE

```
-- Load input from the file named big.txt, and call the single
-- field in the record 'line'.
input = load 'big.txt' as (line);

-- TOKENIZE splits the line into a field for each word.
-- flatten will take the collection of records returned by
-- TOKENIZE and produce a separate record for each one, calling the single
-- field in the record word.
words = foreach input generate flatten(TOKENIZE(line)) as word;

-- Now group them together by each word.
grp = group words by word;

-- Count them.
cntd = foreach grp generate group, COUNT(words);
-- Print out the results.
dump cntd;
```



---

# PIG LATIN - A DATA FLOW LANGUAGE

---

Allows describing:

1. How data flows
2. Should be Read
3. Processed
4. Stored to multiple outputs in parallel

1. Complex Workflows

1. Multiple Inputs are joined
  2. Data is split into multiple streams
2. No ifs and fors



# Map Reduce Programming vs Pig Latin

## PIG LATIN

Provides join, filter, group by, order by, union

Provides optimisation such as rebalancing the reducers

Understand your scripts: Error checking & optimise

Lower cost to write and maintain than Java code for MapReduce.

Brief: 9 lines

Has limitations

## MAP REDUCE

Provides operations such as group by and order by indirectly. Joins are hard

Same optimisations are really hard to write

Your code is opaque to the map reduce

Requires more effort and time

170 lines

More Control.



# Pig Latin vs SQL

## PIG LATIN

User describes exactly how to process the input data.

Long series of data operations

Designed for Hadoop

Targets multiple data operations & parallel processing

Can be used when schema is unknown, incomplete, or inconsistent,

## SQL

Allows users to describe what question they want answered

Around answering one question

Designed for RDBMS


For multiple data operations, the query gets complex

requires well formatted



# PIG - USE CASES

1. ETL
2. Research On Data
3. Iterative processing
  1. Batch processing





---

# PIG PHILOSOPHY

---

## A. Pigs eat anything

- Data: Relational, nested, or unstructured.

## B. Pigs live anywhere

- Parallel processing Language. Not only Hadoop

## C. Pigs are domestic animals

- Controllable, provides custom functions in java/python
- Custom load and store methods

## D. Pigs fly

- Designed for performance



---

# PIG - INSTALLING

---

1. Download 0.13.0 from [Pig's release page](#)
2. `tar -xvfz pig-0.13.0.tar.gz`
3. `mv pig-0.13.0 /usr/local/pig`
4. `export JAVA_HOME=/usr/lib/jvm/java-6-openjdk-amd64`
5. Run `pig -x local` and test if it is working
6. `export PIG_CLASSPATH=/etc/hadoop`
7. Add `pig/bin` folder to `$PATH`
8. Run `pig` and enjoy the grunts



---

# PIG - USING FROM CLOUD LABS

---

- `ssh student@hadoop1.knowbigdata.com` (password is same as mentioned earlier). You can also try `hadoop2`, `3`, `4`
- `pig`
- In HDFS, you have the permissions to manipulate `/users/student/`
- The data has been uploaded to `/users/student/data` in HDFS



---

# PIG - LOCAL MODE

---

- `pig -x local`
- Works without Hadoop installation
- Very useful for testing locally
- Improves the development pace



---

# PIG - COMMAND LINE ARGUMENTS

---

- The first argument can be a pig script file: `pig <myscript>`
- To Execute a single command e.g.:
  - `pig -e ls`
- `-version` , for version
- `-h` see more options



---

# PIG - GRUNT

---

- The prompt “grunt>”
- This is where all of the commands are typed
- You can also control hadoop from here
  - fs (-ls or -cat or any other hadoop command)
  - kill job
  - exec or run pig-script



---

# PIG - BASIC CONSTRUCTION

---


- Each processing step results in new set
- Case Insensitive
- Multiline comments: `/* */`
- Singleline comments: `--`



# PIG - MY FIRST PIG SCRIPT - AVERAGE VAL

Local: data/NYSE\_dividends  
hdfs: /users/student/data/NYSE\_dividends

NYSE	CPO	2009-12-30	0.11
NYSE	CPO	2009-09-28	0.12
NYSE	CPO	2009-06-26	0.13
NYSE	CCS	2009-10-28	0.41
NYSE	CCS	2009-04-29	0.43
NYSE	CIF	2009-12-09	.029
NYSE	CIF	2009-12-09	.028



CPO	0.12
CCS	0.42
CIF	.0285



# PIG - MY FIRST PIG SCRIPT - AVERAGE VAL

Local: data/NYSE\_dividends  
hdfs: /user/student/data/NYSE\_dividends

NYSE	CPO	2009-12-30	0.11
NYSE	CPO	2009-09-28	0.12
NYSE	CPO	2009-06-26	0.13
NYSE	CCS	2009-10-28	0.41
NYSE	CCS	2009-04-29	0.43
NYSE	CIF	2009-12-09	.029
NYSE	CIF	2009-12-09	.028



CPO	0.12
CCS	0.42
CIF	.0285

1. `divs = load '/user/student/data/NYSE_dividends' as (name, ticker, date, val);`
2. `dump divs`
3. `explain divs`
4. `grped = group divs by ticker;`
5. `avged = foreach grped generate group, AVG(divs.val);`
6. `store avged into '/user/student/data-out/avged';`
7. `fs -cat /user/student/data-out/avged/part*`



---

# PIG - DATA TYPES

---

## Scaler

1. int
2. long
3. float - 4 bytes
4. double - 8 bytes
5. chararray
6. bytearray



---

# PIG - DATA TYPES

---

## Complex

### 1. Map

1. ['name'#'bob', 'age'#55]
2. chararray => another complex type or scaler

### 2. Tuple

1. Fixed length, ordered collection
2. Made up fields and their values.
3. Think of it as a row in RDBMS and fields as columns
4. Can refer the fields by position
5. ('bob', 55, 12.3)

### 3. BAG

1. Unordered collection of tuples
2. {('ram', 55, 12.3), ('sally', 52, 11.2)}
  1. Two tuples each with three fields
3. Analogous to SET of rows / dataset
4. Does not need to fit in memory. Spills over to disk



---

# PIG - DATA TYPES

---

## Schemas

1. Not very particular about schema
2. You can either tell upfront, e.g.
  1. `dividends = load 'NYSE_dividends' as (exchange:chararray, symbol:chararray, date:chararray, dividend:float);`
3. Or it will make best type guesses based on data
4. Schema declaration for complex types:
  1. `a:map[]` or `a:map[int] ...`
  2. `a:tuple` or `a:tuple(x, y)` or `a:tuple(x:int, y:bytearray,...)....`
  3. `a:bag{}` or `a:bag{t:(x:int, y:int)}` or `a:bag{t:(x, y)}`
5. You can cast the data from one type to another if possible
  1. similar to Java by affixing “(type)”



---

# MAP / REDUCE

---

**BREAK for 10 mins**





# PIG - LOAD

- **divs = load 'NYSE\_dividends'**
  - Loads the values guessing the datatype
  - Tab is used as separator
  - Absolute or Relative URL would work
- **divs = load 'NYSE\_dividends' using PigStorage(',');**
  - To load comma separated values
- **divs = load 'NYSE\_div\*'**
  - You can specify the wildcard character such as '\*'
- **divs = load 'NYSE\_dividends' using HBaseStorage();**
  - You can load an HBASE table this way
- **divs = load 'NYSE\_dividends' as (name:chararray, ticker:chararray, date, val:float);**
  - Define the data types upfront



---

# PIG - DIAGNOSTIC

---

## **explain dataset**

Shows Detailed analysis of query

## **describe dataset**

Show Schema

## **Syntax Highlighting Tools**

Eclipse <http://code.google.com/p/pig-eclipse>

Emacs <http://github.com/cloudera/piglatin-mode>

TextMate <http://www.github.com/kevinweil/pig.tmbundle>

Vim [http://www.vim.org/scripts/script.php?script\\_id=2186](http://www.vim.org/scripts/script.php?script_id=2186)



---

# PIG - STORE / DUMP

---

## Store

- Stores the data to a file or other storage
- store processed into 'processed' using `PigStorage(',');`

## Dump

- Prints the value on the screen - `print()`
- Useful for
  - debugging or
  - using with scripts that redirect output
- Syntax: *dump variable;* (or relation or field name)



---

# PIG - RELATIONAL OPERATIONS - FOREACH

---

Takes expressions & applies to every record

1. `divs = load '/users/student/data/NYSE_dividends' as (name, ticker, date, val);`
2. `avged = foreach divs generate ticker, val`
3. `dump avged`



# PIG - FOREACH

```
prices = load 'NYSE_daily' as (exchange, symbol, date, open, high, low, close,  
                                volume, adj_close);
```

Basic Operations such as subtraction/addition

- *gain = foreach prices generate close - open;*

Using position references

- *gain2 = foreach prices generate \$6 - \$3;*

Ranges

- For exchange, symbol, date, open
  - *beginning = foreach prices generate ..open;*
- For open, high, low, close
  - *middle = foreach prices generate open..close;*
- For produces volume, adj\_close
  - *end = foreach prices generate volume..;*



# PIG

Will it have any reducer code generated?  
gain = foreach prices generate (close - open);





# PIG

Will it have any reducer code generated?  
gain = foreach prices generate (close - open);



0



# PIG - FOREACH

Extract Data from complex types:

```
bball = load 'baseball' as (name:chararray, team:chararray, traits:bag{}, bat:map[]);  
avg = foreach bball generate bat#'batting_average';
```

## Tuple Projection

```
A = load 'input' as (t:tuple(x:int, y:int));  
B = foreach A generate t.x, t.$1;
```

## Bag Projection

```
A = load 'input' as (b:bag{t:(x:int, y:int)});  
B = foreach A generate b.x;
```

```
A = load 'input' as (b:bag{t:(x:int, y:int)});  
B = foreach A generate b.(x, y);
```



---

# PIG - FOREACH

---

- `divs = load 'NYSE_dividends' as`
- `(exchange:chararray, symbol:chararray, date:chararray, dividends:float);`
- `in_cents = foreach divs generate`
  - `dividends * 100.0 as dividend, dividends * 100.0;`
- `describe in_cents;`

Built-In Function's List

<http://pig.apache.org/docs/r0.13.0/func.html>



# PIG - FILTER

```
-- filter_matches.pig
divs = load 'NYSE_dividends' as
    (exchange:chararray, symbol:chararray, date:chararray, dividends:float);
```

## Regular Expression

1. startswithcm = filter divs by symbol matches 'CM.\*';

## Negation

2. notstartswithcm = filter divs by not symbol matches 'CM.\*';

## Logical Operators

3. result = filter divs by dividends > 0.10 and dividends < .12;

More Operators are are:

<http://pig.apache.org/docs/r0.13.0/basic.html>



Will there be any reducer generated?  
startswithcm = filter divs by symbol matches 'CM.\*';





Will there be any reducer generated?  
startswithcm = filter divs by symbol matches 'CM.\*';



No.



---

# PIG - ORDER BY / DISTINCT

---

## Order BY

```
divs = load 'NYSE_dividends' as  
  (exchange:chararray, symbol:chararray, date:chararray, dividends:float);  
ordered = order divs by symbol, dividends desc;
```

## Distinct

```
divs = load 'NYSE_dividends' as  
  (exchange:chararray, symbol:chararray, date:chararray, dividends:float);  
symbols = foreach divs generate symbol;  
uniqsymbols = distinct symbols;  
dump uniqsymbols;
```



Will it generate any reducer code?  
uniqsymbols = distinct symbols;





Will it generate any reducer code?  
uniqsymbols = distinct symbols;



yes



# PIG - JOIN

## Single Key

```
daily = load 'NYSE_daily' as (exchange, symbol, date, open, high, low, close, volume, adj_close);  
divs = load 'NYSE_dividends' as (exchange, symbol, date, dividends);  
jnd = join daily by symbol, divs by symbol;
```

## Composite Key

```
jnd = join daily by (exchange, symbol), divs by (exchange, symbol);
```

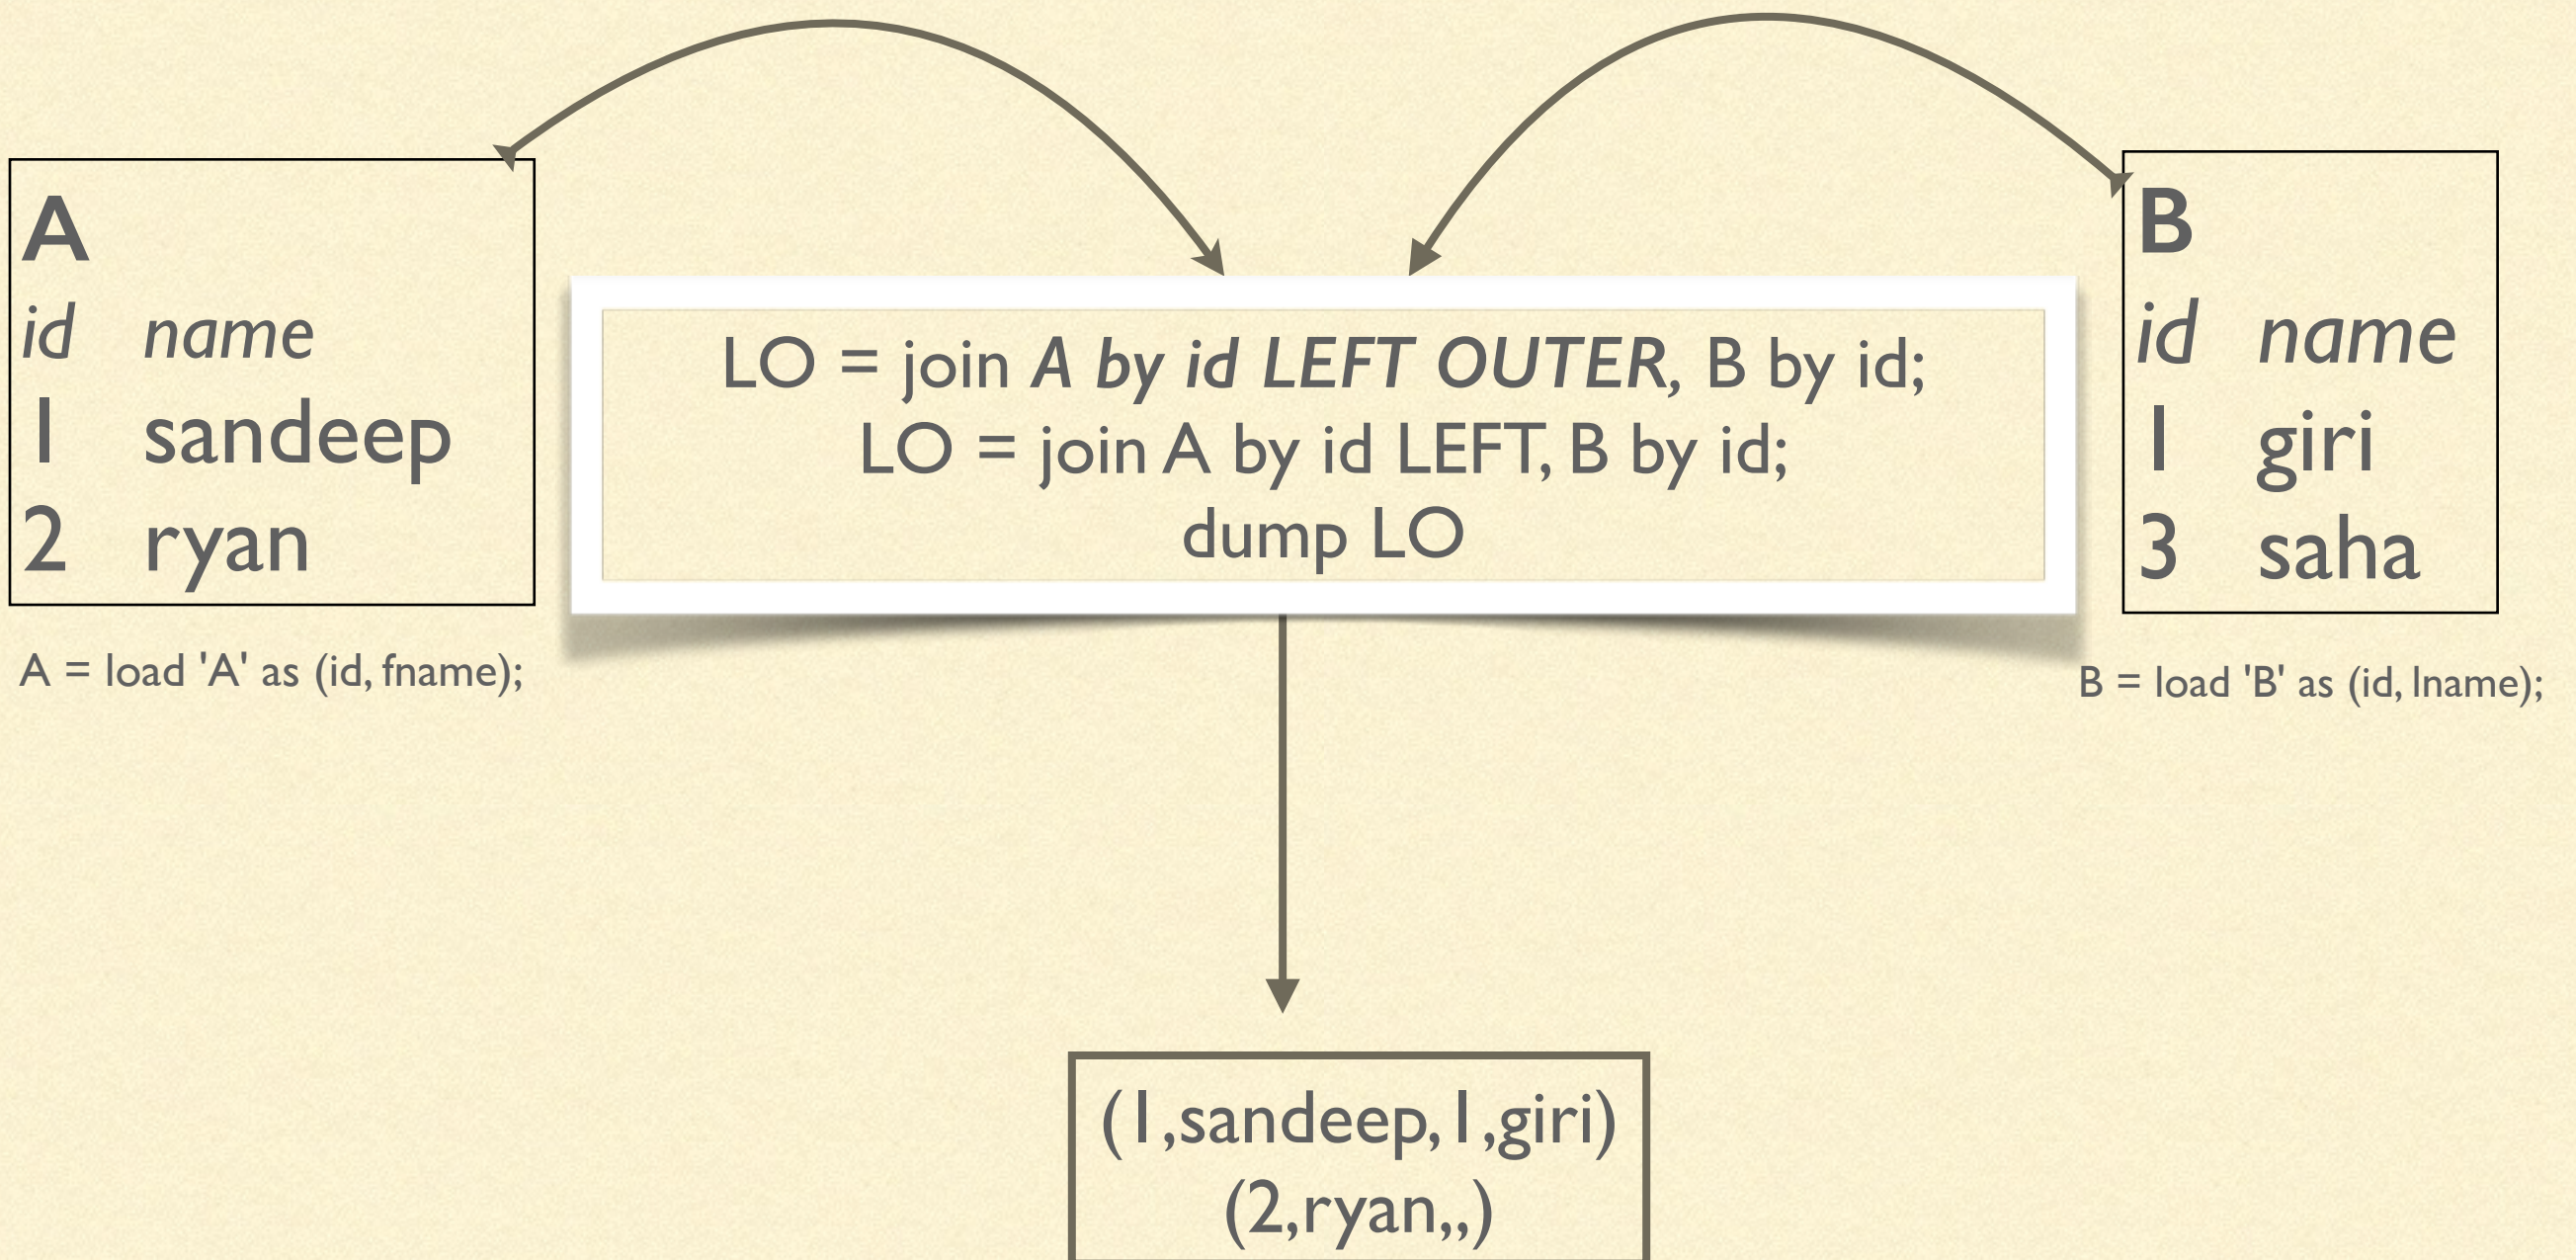
## Self Join - For each stock, find dividends that increased between two dates

```
divs1 = load 'NYSE_dividends' as (exchange:chararray, symbol:chararray, date:chararray, dividends);  
divs2 = load 'NYSE_dividends' as (exchange:chararray, symbol:chararray, date:chararray, dividends);  
jnd = join divs1 by symbol, divs2 by symbol;  
increased = filter jnd by divs1::date < divs2::date and divs1::dividends < divs2::dividends;
```



# PIG - OUTER JOINS - LEFT

Similar to SQL





# PIG - OUTER JOINS - RIGHT

Similar to SQL

A	
id	name
1	sandeep
2	ryan

A = load 'A' as (id, fname);

RO = join A by id **RIGHT OUTER**, B by id;  
or  
RO = join A by id **RIGHT**, B by id;  
dump RO;

B	
id	name
1	giri
3	saha

B = load 'B' as (id, lname);

(1,sandeep,1,giri)  
(,3,saha)





# PIG - OUTER JOINS - FULL

Similar to SQL

A	
id	name
1	sandeep
2	ryan

A = load 'A' as (id, fname);

B	
id	name
1	giri
3	saha

B = load 'B' as (id, lname);


```
RO = join A by id FULL, B by id;  
dump RO;
```

```
(1,sandeep,1,giri)  
(2,ryan,,)  
(,,3,saha)
```






# PIG - OUTER JOINS - QUESTION?



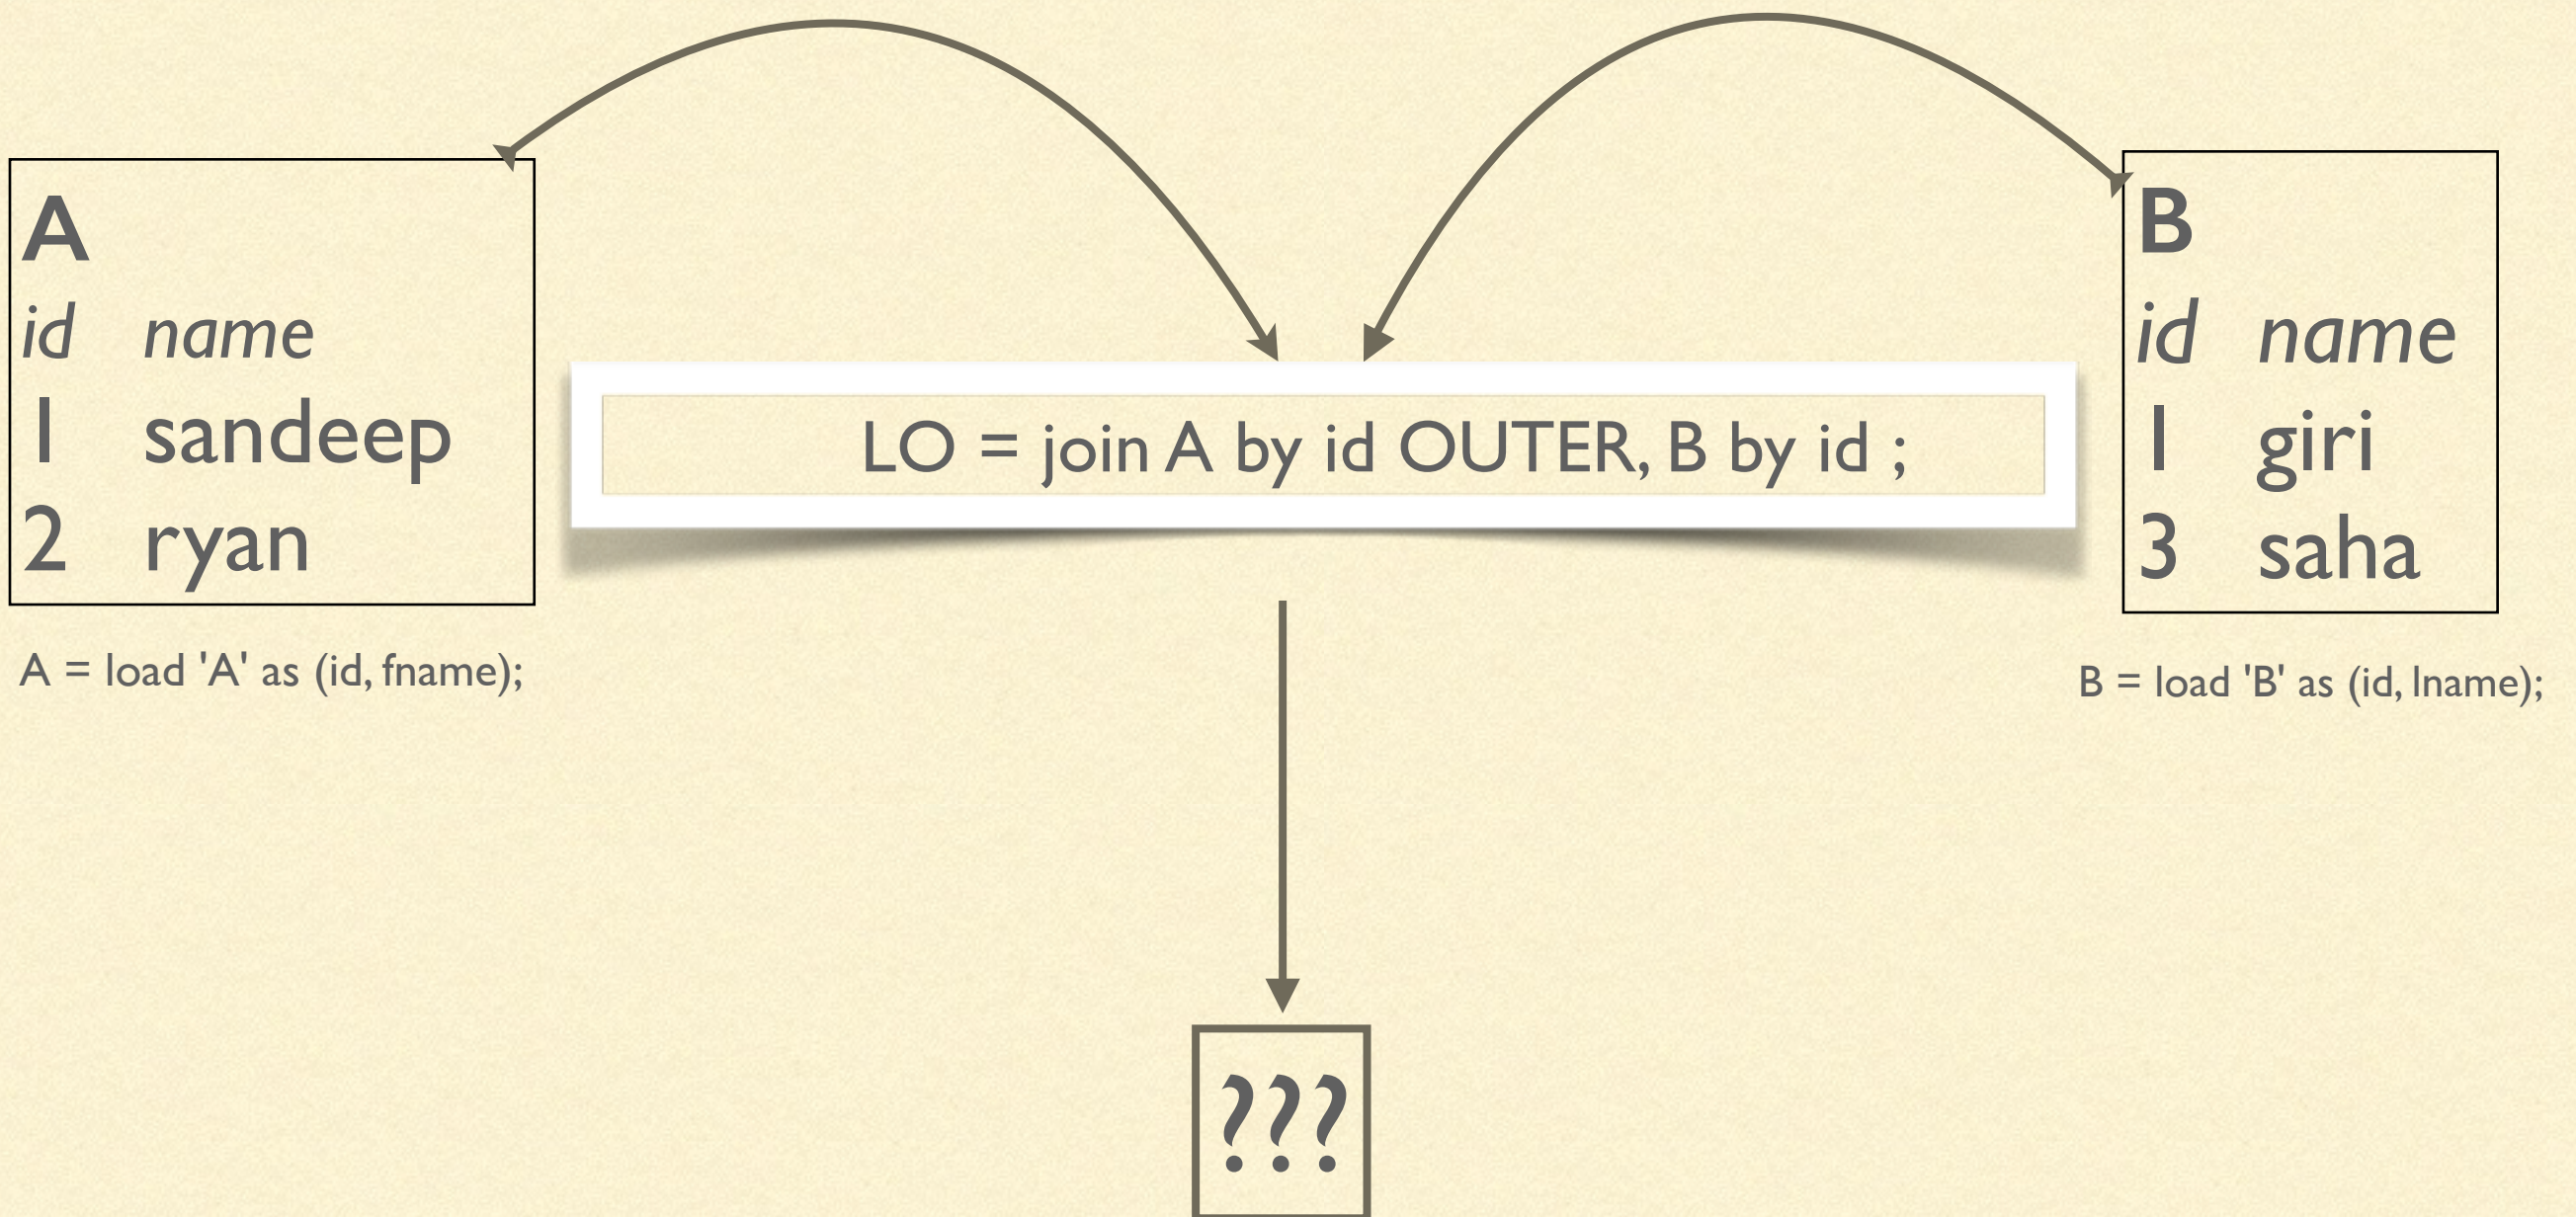


# PIG - OUTER JOINS - QUESTION?



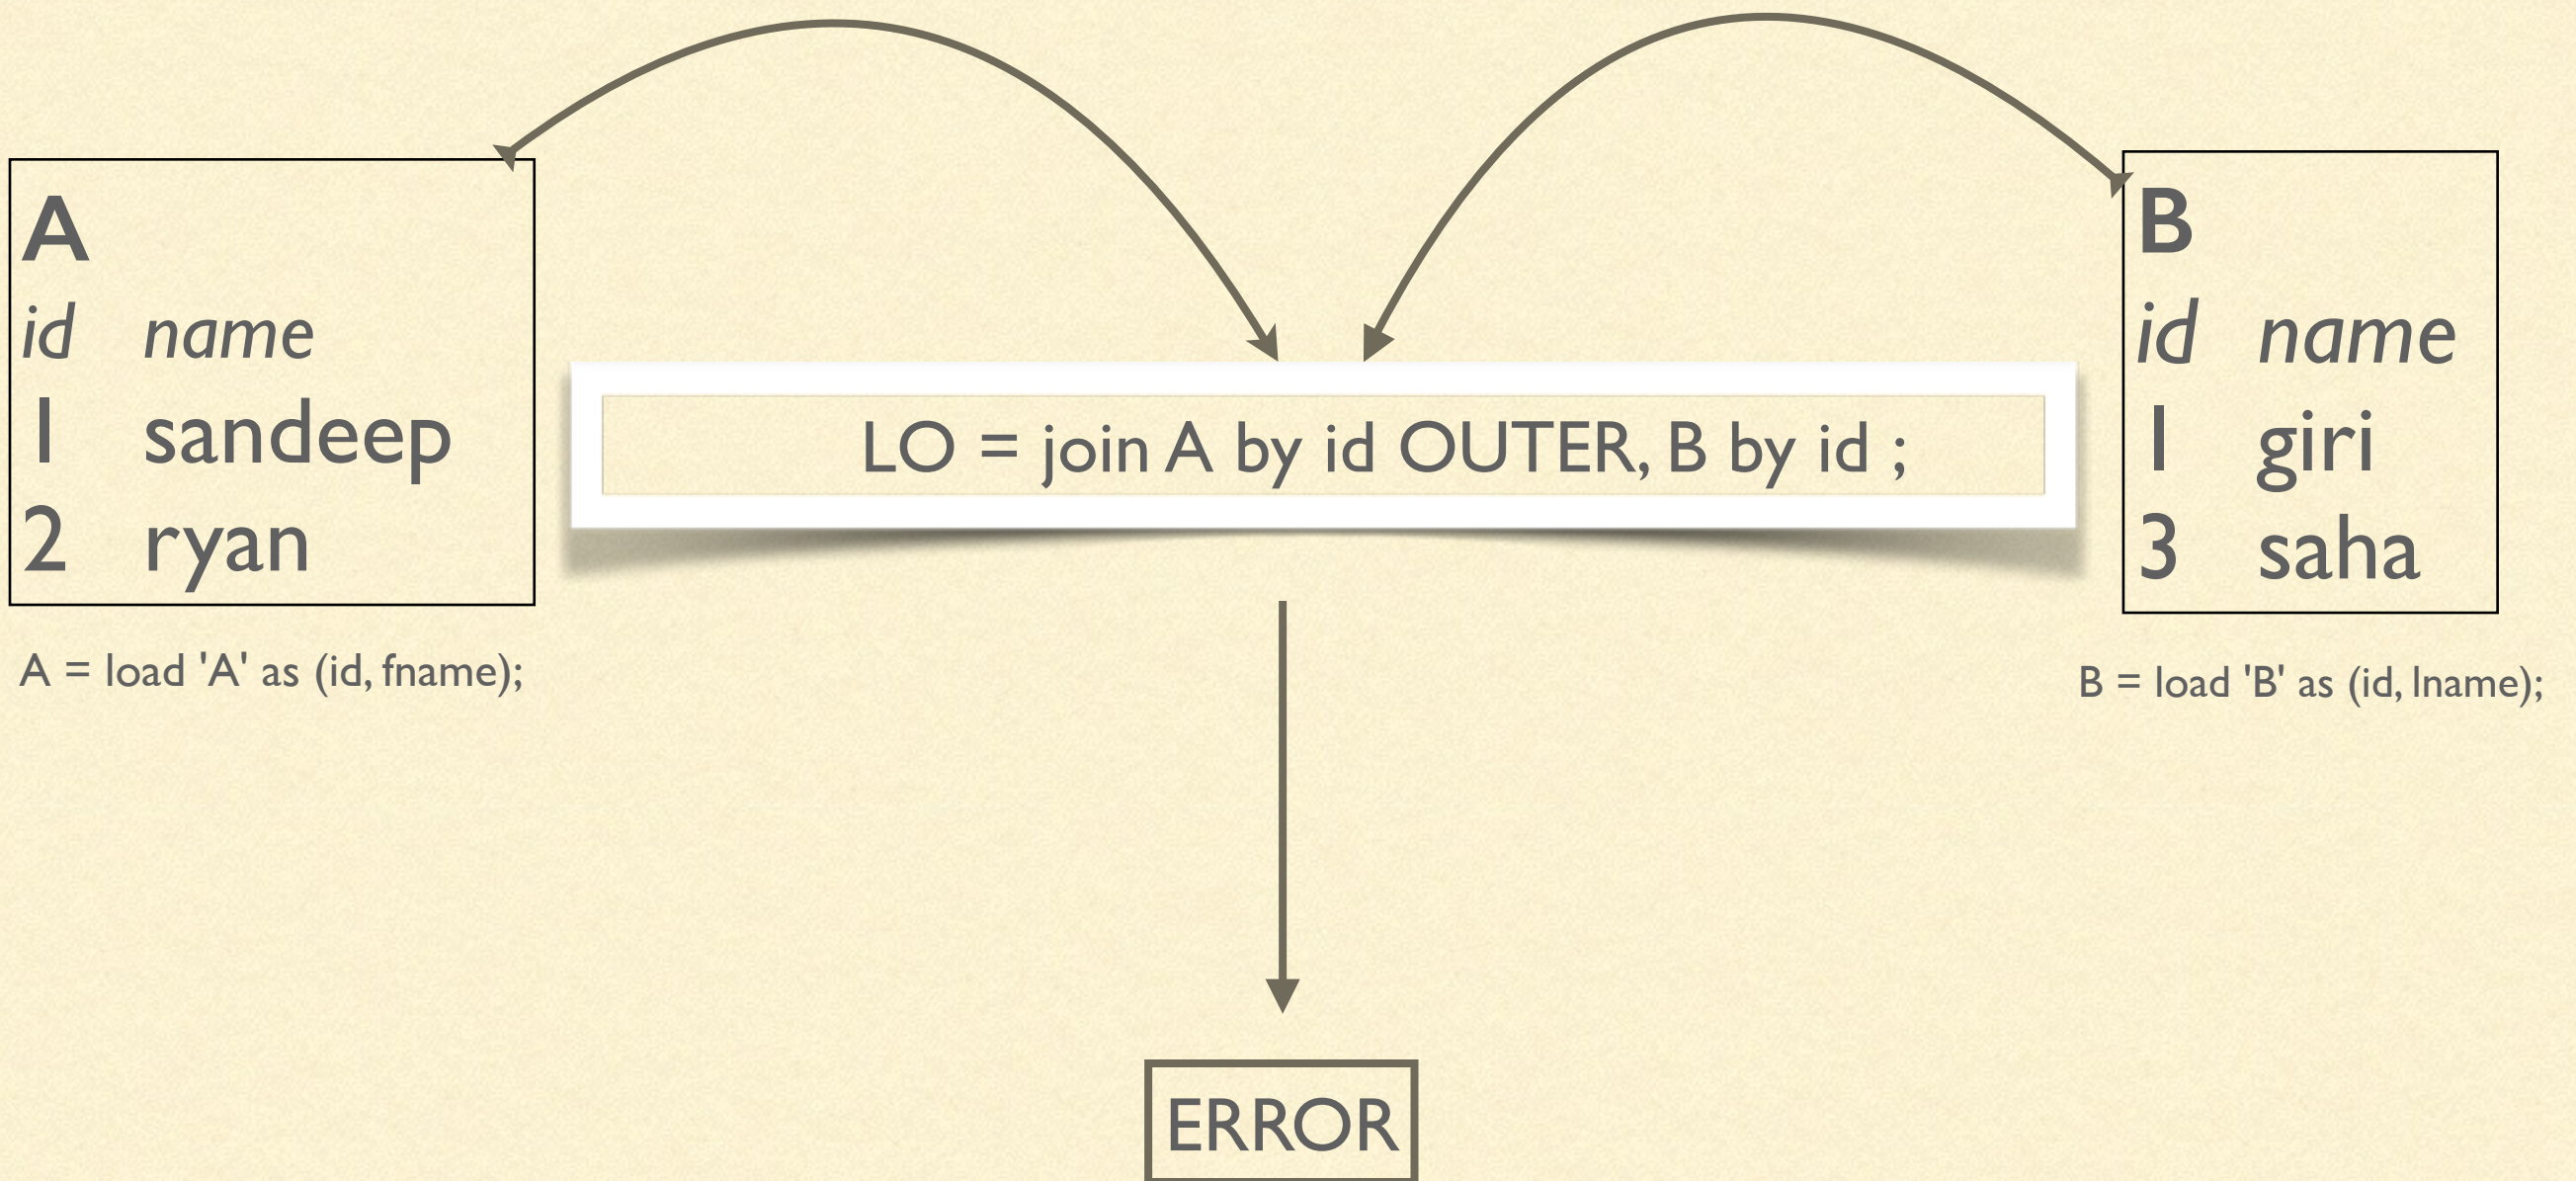


# PIG - OUTER JOINS - QUESTION?





# PIG - OUTER JOINS - QUESTION?





---

# PIG - LIMIT / SAMPLE

---

## LIMIT

```
divs = load 'NYSE_dividends';  
first10 = limit divs 10;  
dump 10;
```

## SAMPLE

```
divs = load 'NYSE_dividends';  
some = sample divs 0.1; --10% records
```



# PIG

```
some = sample divs 0.1;  
dump some;  
dump some;
```

If you run “dump some” twice, will it give the same output?





some = sample divs 0.1;  
If you run “dump some” twice, will it give the same output?



No. Sampling gives randomised output. It is not like limit.



# PIG - PARALLEL

1. Parallelism is the basic promise of pig
2. Basically controls the numbers of reducers
3. Does not impact any mappers

1. `daily = load 'NYSE_daily' as (exchange, symbol, date, open, high, low, close, volume, adj_close);`
2. `bysymb1 = group daily by symbol parallel 10;`
3. — Will create 10 reducers

1. Needs to be applied per statement
2. Script wide: `set default_parallel 10;`
3. Without parallel, allocates reducer per 1 G of input size



# PIG

How many reducers will it produce?  
mylist = foreach divs generate symbol, dividends\*10





How many reducers will it produce?  
mylist = foreach divs generate symbol, dividends\*10



0.



Does parallel clause have any effect in this?  
*mylist = foreach divs generate symbol parallel 10;*  
[Yes/No?]





Does parallel clause have any effect in this?  
*mylist = foreach divs generate symbol parallel 10;*  
[Yes/No?]



N0.



# PIG - FLATTEN

```
sandeep {(c),(c++)}  
arun    {(shell),(c)}  
cmaini  {(java),(c++)}
```

1. *x = load 'myfile' as (name, languages:bag{t:(p)});*
2. *p1 = foreach x generate name, flatten(languages);*
3. *dump p1;*

```
(sandeep,c)  
(sandeep,c++)  
(arun,shell)  
(arun,c)  
(cmaini,java)  
(cmaini,c++)
```



---

# PIG - CALLING EXTERNAL FUNCTIONS

---

## JAVA Built-IN

```
define hex InvokeForString('java.lang.Integer.toHexString', 'int');  
divs = load 'NYSE_daily' as (exchange, symbol, date, open, high, low, close, volume, adj_close);  
inhex = foreach divs generate symbol, hex((int)volume);
```

## Your Own Java Function

```
register 'acme.jar';  
define convert com.acme.financial.CurrencyConverter('dollar', 'euro');  
divs = load 'NYSE_dividends' as (exch:chararray, sym:chararray, date:chararray, vals:float);  
backwards = foreach divs generate convert(dividends);
```

## Calling a Python function

```
register 'production.py' using jython as bball;  
calcprod = foreach nonnull generate name, bball.production(somval);
```



# PIG - FOREACH NESTED

```
blore sandeep  
blore sandeep  
blore vivek  
nyc john  
nyc john
```

Find count of distinct names in each city

???

```
(blore,2)  
(nyc,1)
```



# PIG - FOREACH NESTED

```
blore sandeep  
blore sandeep  
blore vivek  
nyc john  
nyc john
```

Find total distinct names in each city

```
1. people = load 'people' as (city, name);  
2. grp = group people by city;  
3. counts = foreach grp {  
    1. dnames = distinct people.name;  
    2. generate group, COUNT(dnames);  
4. }
```

```
(blore,2)  
(nyc,1)
```



# PIG - FRAGMENTED JOINS

Translate the countries to country code

NAME	COUNTRY
------	---------

sandeep IN

kumar US

gerald CN

... 1 billion entries ...

COUNTRY	CODE
---------	------

IN 91

US 1

CN 86

... 100 entries ...

NAME	CODE
------	------

sandeep 91

kumar 1

gerald 86

... 1 billion entries ...

/.

????



# PIG - FRAGMENTED JOINS - REPLICATED

Translate the countries to country code

NAME	COUNTRY
------	---------

sandeep IN

kumar US

gerald CN

... 1 billion entries ...

COUNTRY	CODE
---------	------

IN 91

US 1

CN 86

... 100 entries ...

NAME	CODE
------	------

sandeep 91

kumar 1

gerald 86

... 1 billion entries ...

1. *persons = load 'persons' as (name, country);*
2. *cc = load 'cc' as (country, code);*
3. *joined = join persons by country, cc by country using 'replicated';*



---

# PIG - JOINING SKEWED DATA

---

Skewed data:

- Too many values for some keys

- Too few keys for some

```
users = load 'users' as (name:chararray, city:chararray);  
cinfo = load 'cityinfo' as (city:chararray, population:int);  
jnd   = join cinfo by city, users by city using 'skewed';
```

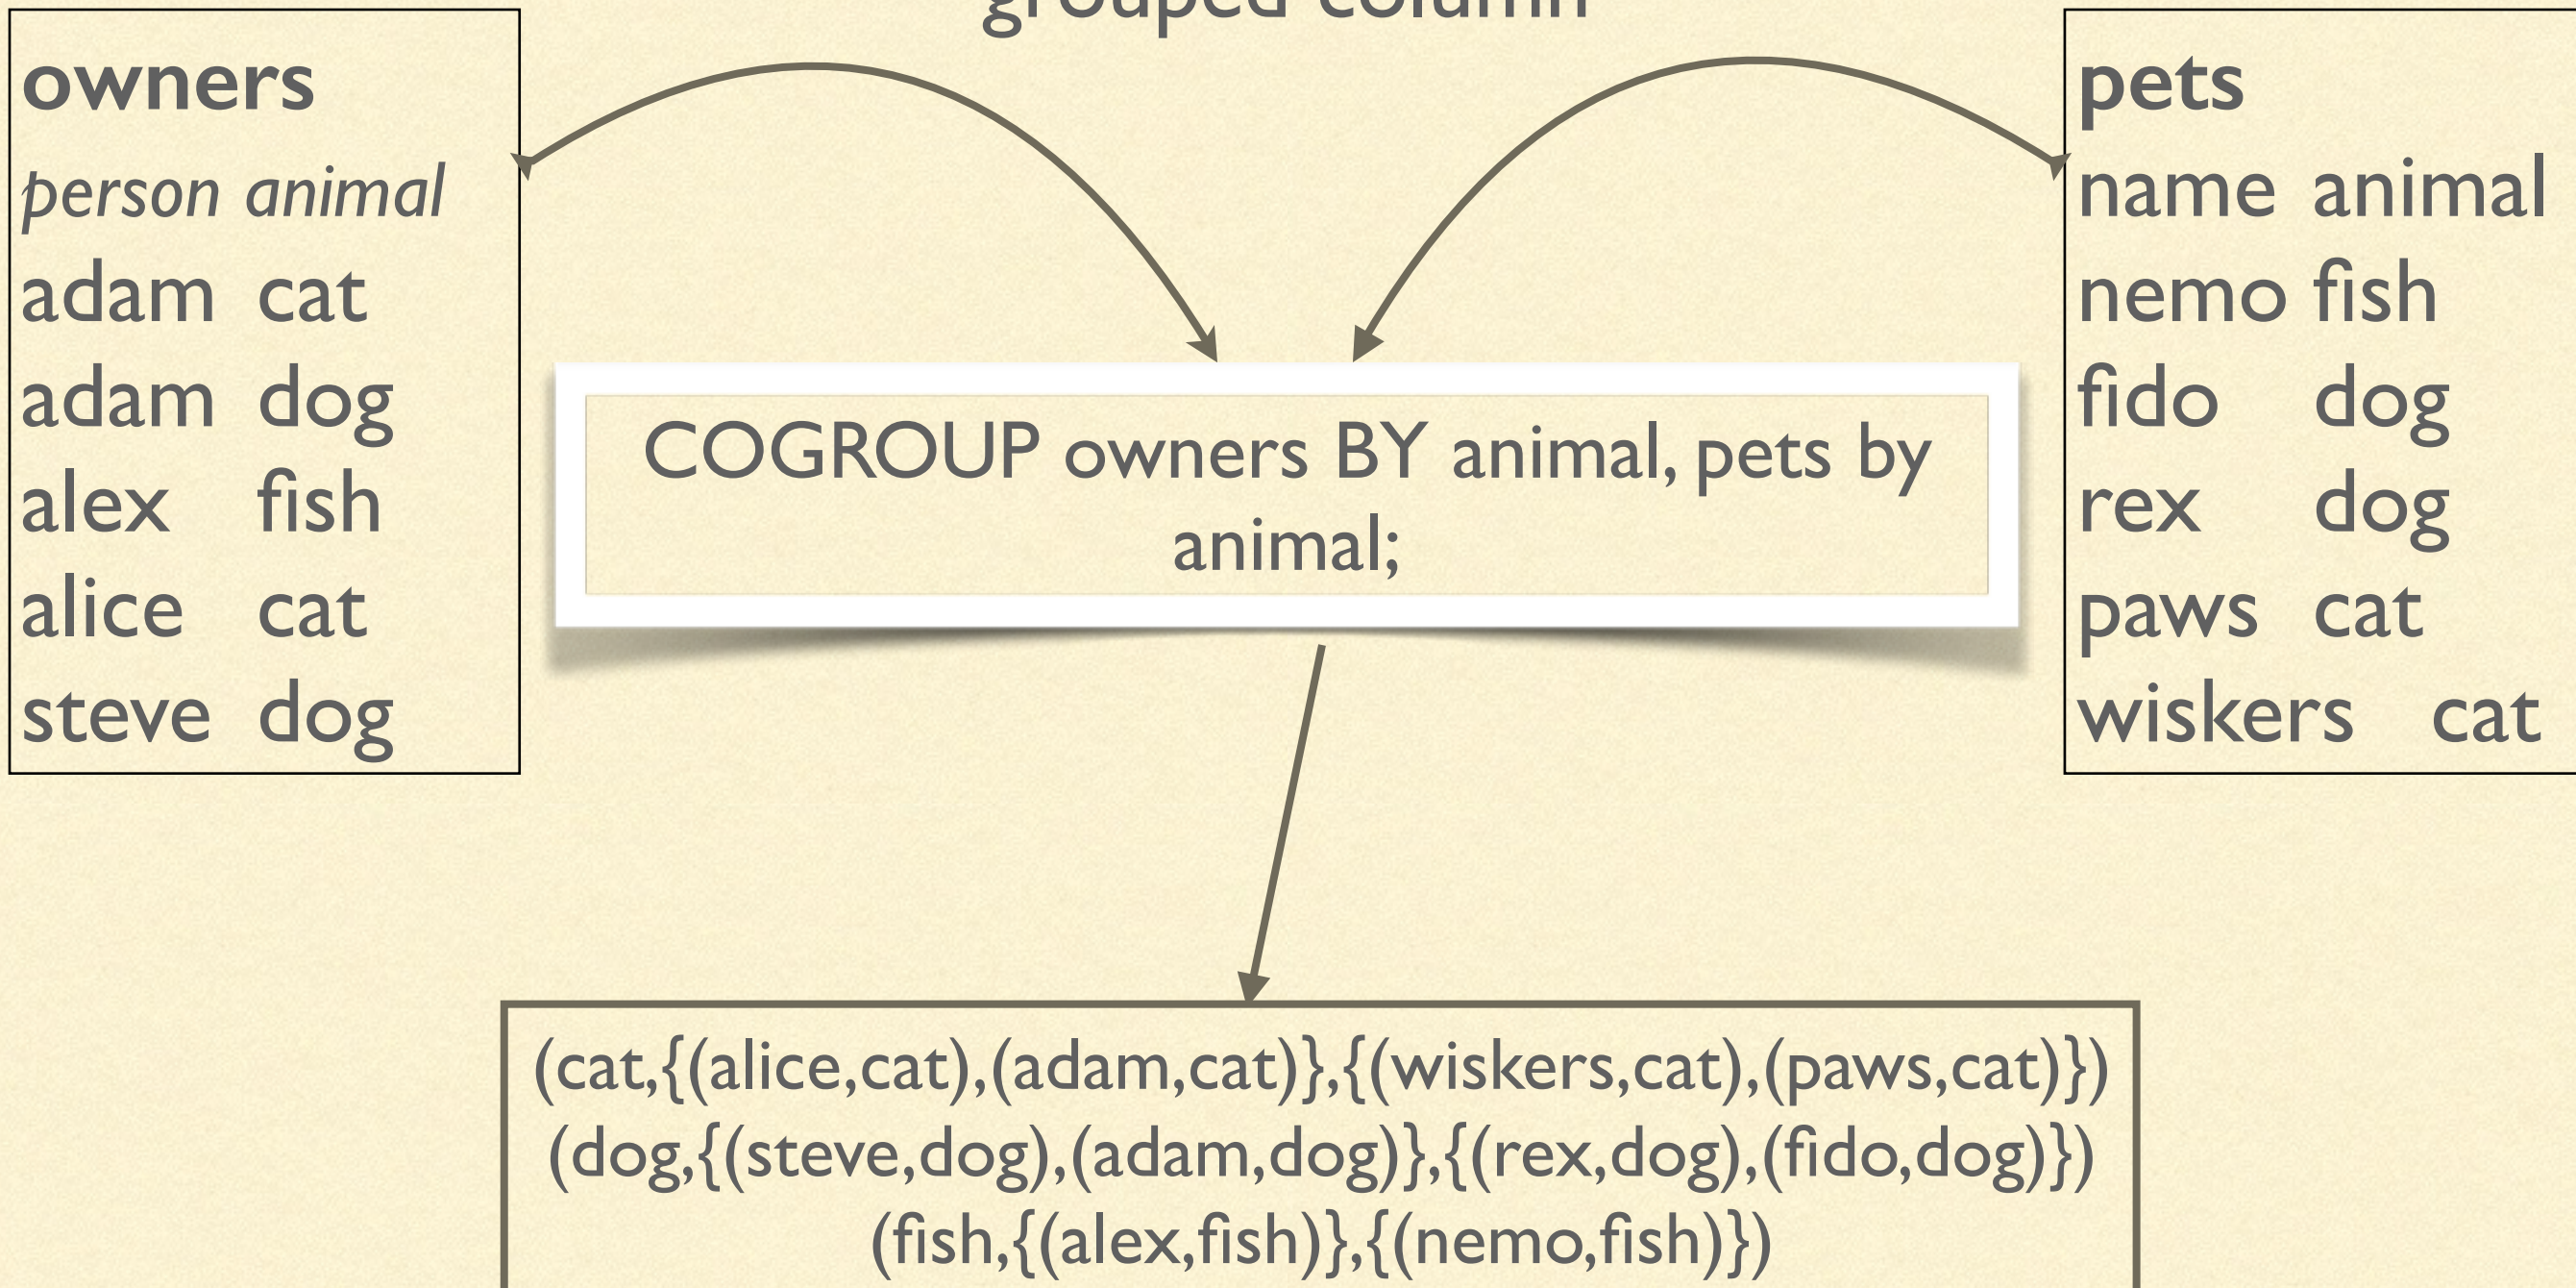
# PIG - JOINING SORTED DATA

```
jnd   = join daily by symbol, divs by symbol using 'merge';
```



# PIG - COGROUP

Group two tables by a column and then join on the grouped column





# PIG - UNION

Concatenating two data sets.  
**Unix Equivalent: cat file1 file2**  
**SQL: union all**

```
A = load '/user/me/data/files/input1';  
B = load '/user/someoneelse/info/input2';  
C = union A, B;
```

**TO force schema use**

```
A = load '/user/me/data/files/input1' as (x, y);  
B = load '/user/someoneelse/info/input2' (x,z);  
C = union onschema A, B; -- x,y, z as columns
```



---

# PIG - CROSS

---

Produce Cross Product

```
tonsodata = cross daily, divs parallel 10;
```



# PIG - STREAM - INTEGRATE WITH LEGACY CODE

Say you have a propriety or legacy code that need to applied.

- Invoked once on every map or reduce task
- Not on every record

```
divs = load 'NYSE_dividends' as  
      (exchange, symbol, date, dividends);  
highdivs = stream divs through `highdiv.pl` as  
           (exchange, symbol, date, dividends);
```

**The program is shipped to every task node:**

```
define hd `highdiv.pl -r xyz` ship('highdiv.pl');  
divs = load 'NYSE_dividends' as (exchange, symbol, date, dividends);  
highdivs = stream divs through hd as (exchange, symbol, date, dividends);
```



---

# PIG - PARAMETER SUBSTITUTION

---

Like with any scripting, it is possible to pass an argument  
And use it in the script

```
--daily.pig  
...  
yesterday = filter daily by date == '$DATE';
```

```
pig -p DATE=2009-12-17 daily.pig
```

```
#Param file daily.params  
YEAR=2009-  
MONTH=12-  
DAY=17  
DATE=$YEAR$MONTH$DAY
```

```
pig -param_file daily.params daily.pig
```



---

## PIG - SPLITS

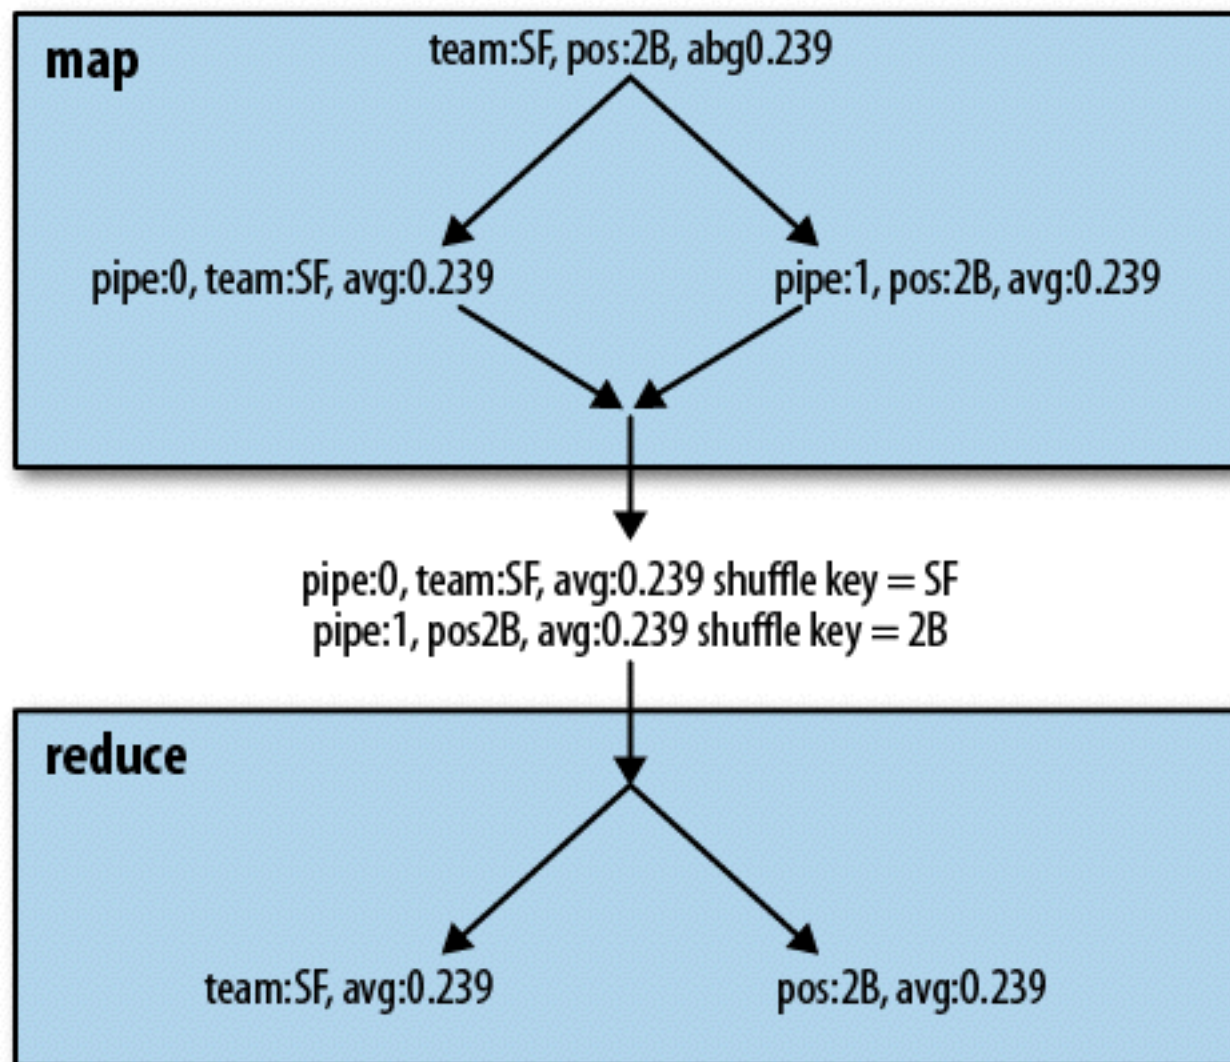
---

```
wlogs = load 'weblogs' as (pageid, url, timestamp);  
split wlogs into  
    apr03 if timestamp < '20110404',  
    apr02 if timestamp < '20110403' and timestamp > '20110401',  
    apr01 if timestamp < '20110402' and timestamp > '20110331';  
store apr03 into '20110403';  
store apr02 into '20110402';  
store apr01 into '20110401';
```

1. A single record can go to multiple legs
2. A record could also be dropped out



# PIG - NON LINEAR EXECUTION



Pig tries to combine multiple groups and filters into one.  
And also creates multiple map jobs depending on what we are trying to do.



---

# PIG - REFERENCES

---

General Docs:

<http://pig.apache.org/docs/r0.13.0/>

Creating Custom Load Functions:

<http://pig.apache.org/docs/r0.13.0/udf.html>





# Big Data & Hadoop

---

Thank you.



**+91-9538998962**

**sandeep@knowbigdata.com**